

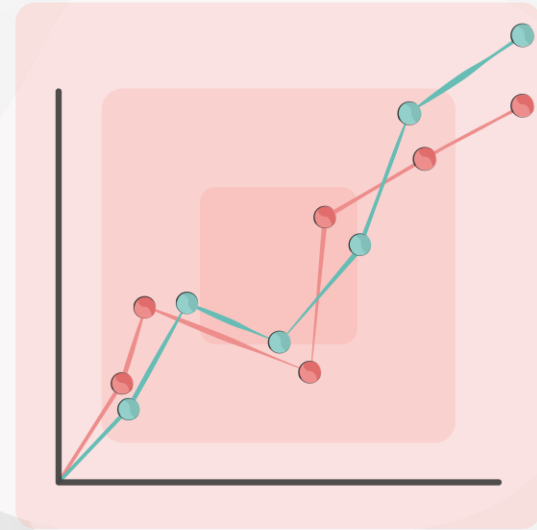


Level 5 Data Engineer

Module 5 Topic 3

Containers and Orchestration

**Welcome to today's
webinar.**


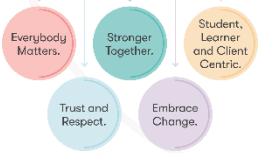


Session objectives

This webinar supports the following learning outcomes:

- Discuss how cloud automation and orchestration can help streamline the management and maintenance of cloud-based systems and applications;
- Analyse how the use of tools and technologies such as Ansible, Puppet, and Chef may be used to automate routine tasks in order to reduce the potential for human error in cloud-based systems and applications;
- Use tools and technologies such as Kubernetes, Docker Swarm, and Mesos to manage containers and balance workloads to ensure they are working together effectively and high availability exists.

Building Careers
Through Education



Container basics

Containers are a method of operating system virtualisation


Benefits –

- Repeatable.
- Self-contained environments.
- Software runs the same in different environments.
 - Developer's laptop, test, production.
- Faster to launch and stop or terminate than virtual machines

Your Container



Building Careers
Through Education

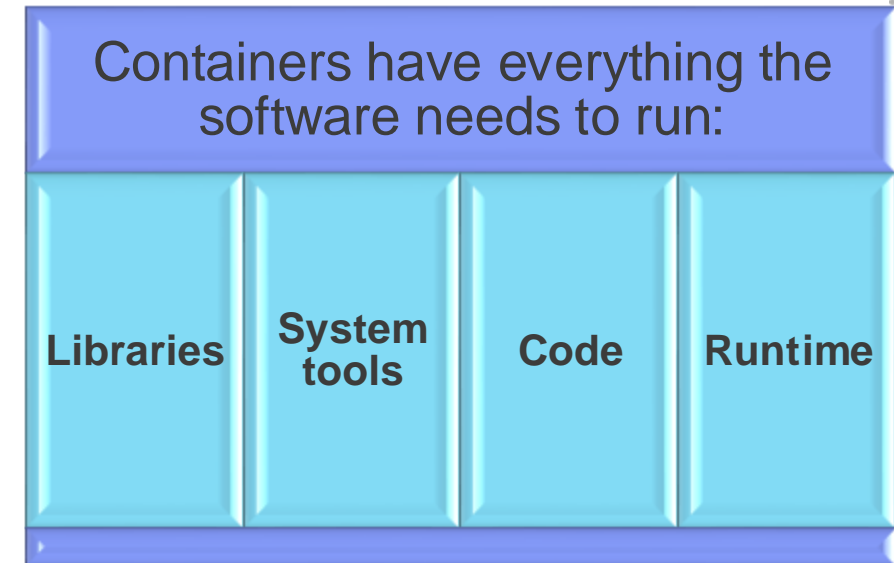


What is Docker?

- **Docker** is a software platform that enables you to build, test, and deploy applications quickly.
- You run containers on Docker.
 - Containers are created from a template called an *image*.
- A **container** has everything a software application needs to run.



Container



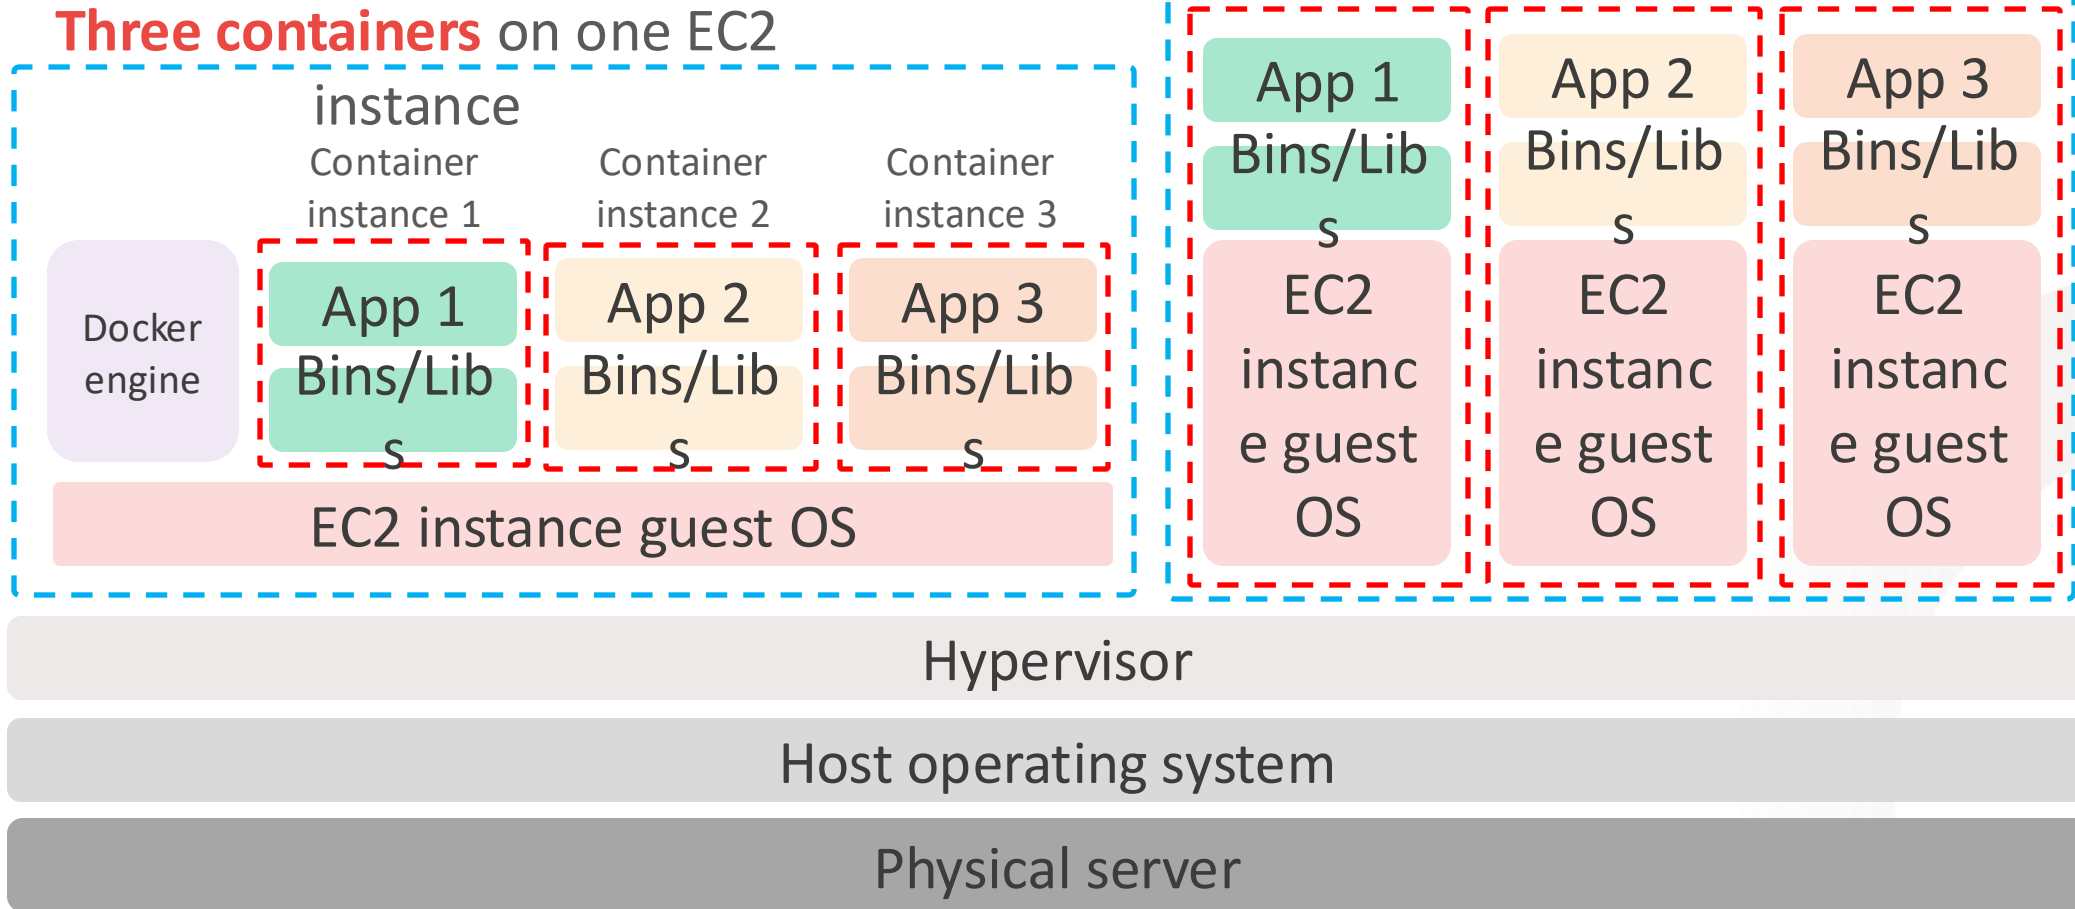
Containers versus virtual machines

Building Careers
Through Education



Example

Three virtual machines on three EC2 instances




Part of
AWS Global
Infrastructure



Amazon Elastic Container Service (ECS)

- Amazon Elastic Container Service (**Amazon ECS**) –
 - A highly scalable, fast, **container management service**
- Key benefits:
 - Orchestrates the running of Docker containers
 - Maintains and scales the fleet of nodes that run your containers
 - Removes the complexity of standing up the infrastructure
- Integrated with features that are familiar to Amazon EC2 service users:
 - Elastic Load Balancing
 - Amazon EC2 security groups
 - Amazon EBS volumes
 - IAM roles


Building Careers
Through Education




**Amazon Elastic
Container Service**



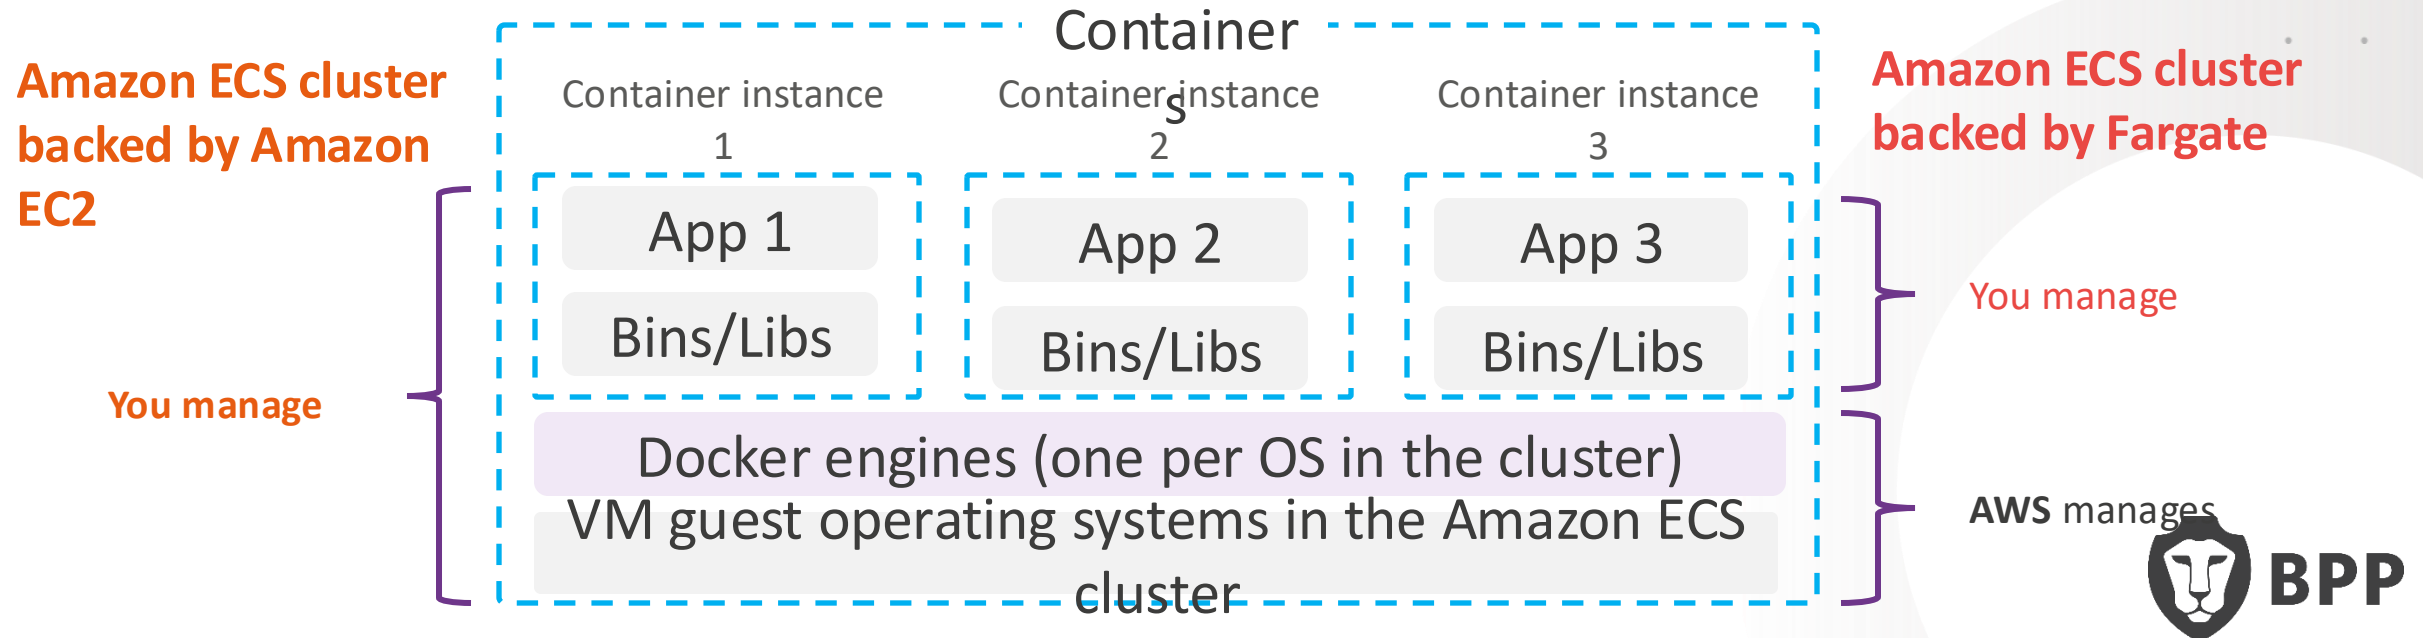
Amazon ECS orchestrates containers



Amazon ECS cluster options




- **Key question:** Do **you** want to manage the Amazon ECS cluster that runs the containers?
- If **yes**, create an **Amazon ECS cluster backed by Amazon EC2** (provides more granular control over infrastructure)
- If **no**, create an **Amazon ECS cluster backed by AWS Fargate** (easier to maintain, focus on your applications)



What is Kubernetes?

- **Kubernetes is open source software for container orchestration.**
 - Deploy and **manage containerised applications** *at scale*.
 - The same toolset can be used on premises and in the cloud.
- **Complements Docker.**
 - Docker enables you to run multiple containers on a single OS host.
 - Kubernetes **orchestrates** multiple Docker hosts (nodes).
- **Automates –**
 - Container provisioning.
 - Networking.
 - Load distribution.
 - Scaling.

Building Careers
Through Education




Activity

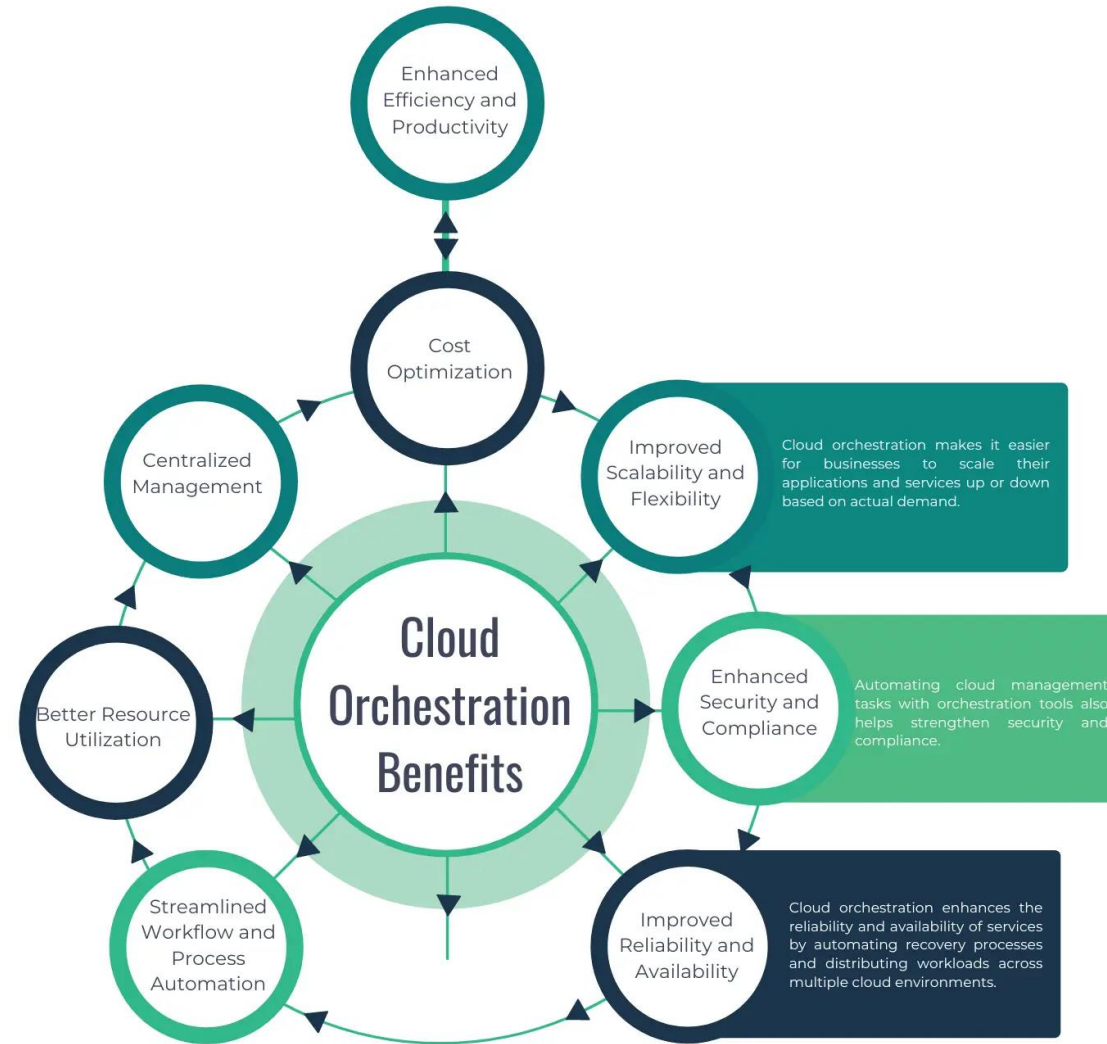
We are now going to:

1. Discuss how cloud automation and orchestration can help streamline the management and maintenance of cloud-based systems and applications;
2. Analyse how the use of tools and technologies such as Ansible, Puppet, and Chef may be used to automate routine tasks in order to reduce the potential for human error in cloud-based systems and applications;
3. Use tools and technologies such as Kubernetes, Docker Swarm, and Mesos to manage containers and balance workloads to ensure they are working together effectively and high availability exists.
4. Demonstration of Kubernetes on Azure

Building Careers
Through Education

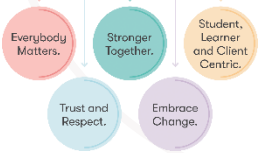


Introduction to Cloud orchestration



- **Source:** <https://www.mega.com/blog/what-is-cloud-orchestration>

Building Careers
Through Education



© MEGA International




Discussion

Use cases:

- Automated Deployment and Scaling of Web Applications
- Continuous Integration and Continuous Deployment (CI/CD)
- Infrastructure as Code (IaC)
- Disaster Recovery and Backup Automation
- Automated Security Compliance and Patching
- Multi-Cloud Management
- Automated Compliance Auditing

Building Careers
Through Education



Automation tools – Chef, Puppet & Ansible

Ansible:

- Agentless architecture
- Uses YAML for defining playbooks
- Easy to learn and use

Puppet:

- Declarative language for defining system configurations
- Uses a client-server model
- Suitable for managing large infrastructures

Chef:

- Uses Ruby-based DSL for writing configuration scripts (cookbooks)
- Follows a client-server architecture
- Focuses on infrastructure as code

Building Careers
Through Education



Ansible example – Deploy a webserver

```
---
- name: Deploy a web server
  hosts: webserver
  become: yes

  vars:
    apache_packages:
      - apache2
      - apache2-utils

  tasks:
    - name: Update and upgrade apt packages
      apt:
        update_cache: yes
        upgrade: dist


    - name: Install Apache packages
      apt:
        name: "{{ apache_packages }}"
        state: present

    - name: Start and enable Apache service
      service:
        name: apache2
        state: started
        enabled: yes

    - name: Create a custom index.html
      copy:
        dest: /var/www/html/index.html
        content: |
          <!DOCTYPE html>
          <html>
          <head>
            <title>Welcome to your web server!</title>
          </head>
          <body>
            <h1>Hello, World!</h1>
            <p>This is a custom web server deployed with Ansible.</p>
          </body>
          </html>
      notify:
        - Restart Apache

  handlers:
    - name: Restart Apache
      service:
        name: apache2
        state: restarted
```

Building Careers
Through Education



Puppet manifest – deploy a webserver

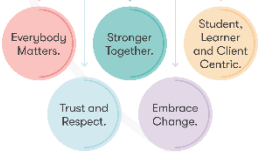
```
# webserver.pp

node 'webserver' {
  # Ensure Apache is installed
  package { 'apache2':
    ensure => present,
  }

  # Ensure the Apache service is running and enabled at boot
  service { 'apache2':
    ensure => running,
    enable => true,
  }

  # Create a custom index.html file
  file { ['/var/www/html/index.html']:
    ensure => file,
    content => '<html>
      <head>
        <title>Welcome to your web server!</title>
      </head>
      <body>
        <h1>Hello, World!</h1>
        <p>This is a custom web server deployed with Puppet.</p>
      </body>
    </html>',
    mode    => '0644',
    owner   => 'www-data',
    group   => 'www-data',
    require => Package['apache2'],
  }
}
```

Building Careers
Through Education



Chef recipe – deploy a webserver

```
# webserver.rb

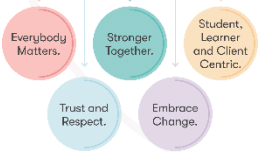
# Update the package database
apt_update 'update' do
  action :update
end

# Install Apache package
package 'apache2' do
  action :install
end



# Enable and start Apache service
service 'apache2' do
  action [:enable, :start]
end

# Create a custom index.html file
file '/var/www/html/index.html' do
  content '<html>
    <head>
      <title>Welcome to your web server!</title>
    </head>
    <body>
      <h1>Hello, World!</h1>
      <p>This is a custom web server deployed with Chef.</p>
    </body>
  </html>'
  mode '0644'
  owner 'www-data'
  group 'www-data'
end
```

Building Careers
Through Education



Ansible, Puppet & Chef - Differences

Category	 Chef	 Puppet	 Ansible
Availability (in case of failure)	Backup Server	Alternative Master	Secondary instance
Configuration Language	Ruby DSL	Ruby, Puppet DSL, Embedded Ruby (ERB), DSL	Python, YAML
Architecture	Master-Agent	Master-Agent	Only Master (Agentless)
Installation Process	Time-intensive and complex due to Chef Workstation	Time-intensive due to master-agent certificate signing	Easy
Configuration Management	Pull	Pull	Push and Pull
Scalability	High	High	Very High
Interoperability	Chef Server works only on Linux/Unix; Chef Client and Workstation can work on Windows as well	Puppet Master works only on Linux/Unix; Puppet Agent or Client works on Windows	Ansible Server works on Linux/Unix; Client machines on Windows
Capabilities	<ul style="list-style-type: none"> Continuous delivery with automated workflow Compliance and security management Infrastructure automation 	<ul style="list-style-type: none"> Orchestration Automated provisioning Code and node management Configuration automation Visualization and reporting High transparency Role-based access control 	<ul style="list-style-type: none"> Simple orchestration Streamlined provisioning Continuous delivery with automated workflow App deployment Security and compliance integration into automation
Pricing	<ul style="list-style-type: none"> Standard Hosted Chef - USD 72/year/node Chef Automation - USD 137/year/node 	<ul style="list-style-type: none"> Puppet Enterprise - USD 120/year/node Premium - USD 199/year/node 	<ul style="list-style-type: none"> Self-support Package - USD 5000/year Premium - USD 14,000/year/100 nodes

veritis
transcend




1-877-VERITIS (283-7484), 972-753-0022



connect@veritis.com

Building Careers
Through Education




- Source: <https://www.veritis.com/blog/chef-vs-puppet-vs-ansible-comparison-of-devops-management-tools/>

Best practices for automation tools

1. Modular Design
2. Version Control
3. Idempotency
4. Testing
5. Documentation
6. Security
7. Consistency
8. Monitoring and Logging
9. Collaboration
10. Continuous Integration/Continuous Deployment (CI/CD)

Building Careers
Through Education



Managing and balancing workloads

Key Objectives:


- Automate Deployment: Simplify the process of deploying applications.
- Scale Applications: Dynamically adjust the number of running containers based on demand.
- Manage Resources: Optimise the allocation and use of computing resources.

Benefits of Using Orchestration Tools:

- High Availability: Ensure that applications are always available, even during failures.
- Load Balancing: Distribute traffic evenly across multiple containers to prevent overload.
- Self-Healing: Automatically detect and replace failed containers.
- Flexibility: Support a wide range of applications and workloads.

Popular Orchestration Tools:

- Kubernetes: Advanced orchestration features, suitable for complex, large-scale applications.
- Docker Swarm: Simple setup and management, integrates seamlessly with Docker.
- Apache Mesos: Suitable for large-scale distributed systems and mixed workloads.



Kubernetes overview

Open-source platform for automating deployment, scaling, and operations of application containers

Manages clusters of containers

Key Features:

- **Self-Healing:** Automatically replaces failed containers
- **Horizontal Scaling:** Scales applications up or down based on demand
- **Service Discovery and Load Balancing:** Distributes traffic across containers

Use Cases:

- Managing microservices architectures
- Scaling web applications

Building Careers
Through Education



Docker Swarm Overview

Native clustering and orchestration solution for Docker

Simple setup and integration with Docker tools


Key Features:

- **Declarative Service Model:** Define the desired state of services
- **Multi-Host Networking:** Connects containers across multiple hosts
- **Load Balancing:** Distributes incoming requests across container instances

Use Cases:

- Small to medium-sized deployments
- Simple, straightforward container orchestration

Building Careers
Through Education



Apache Mesos Overview

Cluster manager for large-scale distributed systems

Supports various container formats and workloads

Key Features:

- **Resource Isolation:** Allocates resources dynamically based on needs
- **Multi-Framework Support:** Runs different frameworks (e.g., Marathon, Spark)
- **High Availability:** Ensures continuous operation even if some nodes fail

Use Cases:

- Large-scale data processing
- Mixed workloads and environments

Building Careers
Through Education



Discussion

AWS and orchestration

1. How does AWS handle orchestration?

Building Careers
Through Education




**Submit your responses to
the chat!**



Amazon Elastic Kubernetes Service (EKS)

- Amazon Elastic Kubernetes Service (**Amazon EKS**)
 - Enables you to run Kubernetes on AWS
 - Certified Kubernetes conformant (supports easy migration)
 - Supports Linux and Windows containers
 - Compatible with Kubernetes community tools and supports popular Kubernetes add-ons
- Use Amazon EKS to –
 - Manage clusters of Amazon EC2 compute instances
 - Run containers that are orchestrated by Kubernetes on those instances



**Amazon Elastic
Kubernetes Service**



Amazon Elastic Container Registry (Amazon ECR)



Amazon ECR is a fully managed Docker **container registry** that makes it easy for developers to store, manage, and deploy Docker container images.



Amazon Elastic
Container Registry



Image Registry




Key takeaways

The key takeaways from this section are as follows:

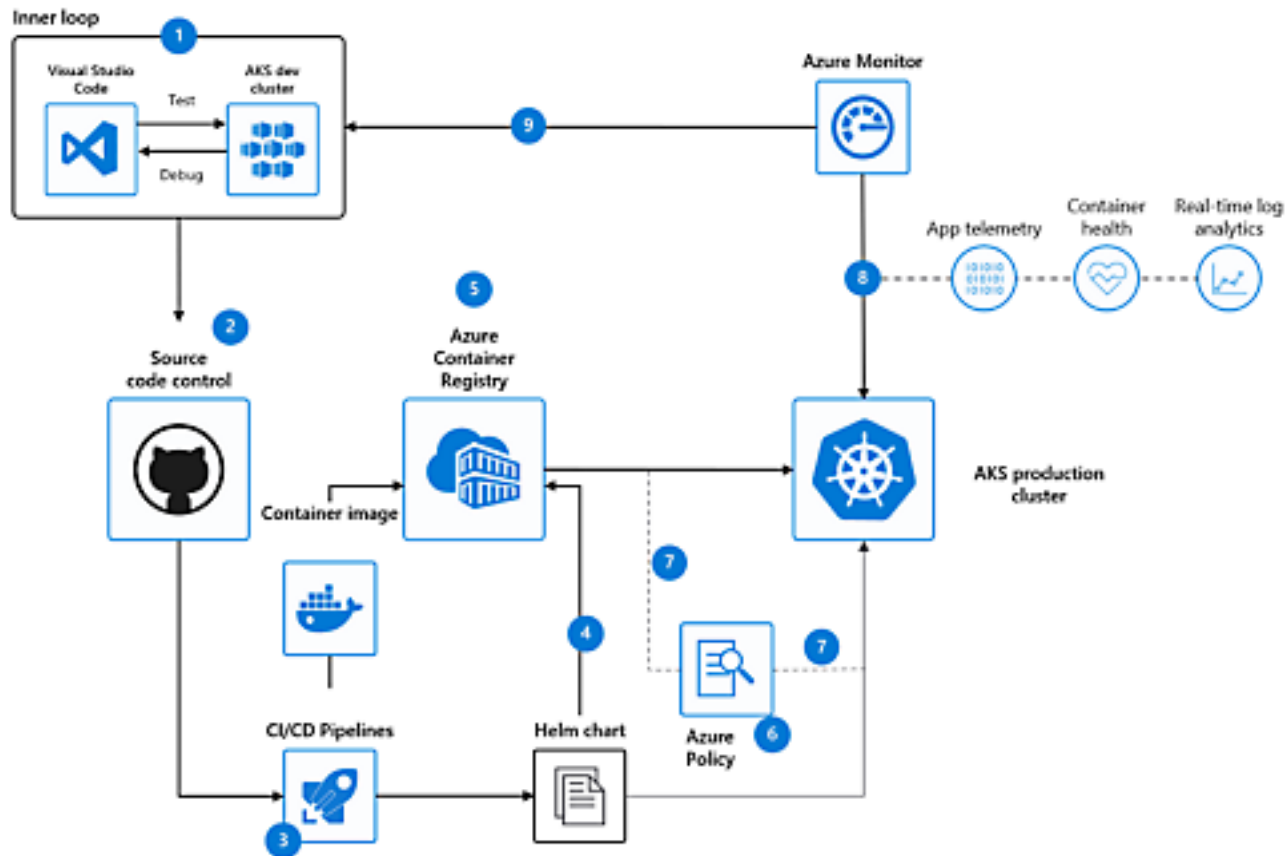
- Containers can hold everything that an application needs to run.
- Docker is a software platform that packages software into containers.
- A single application can span multiple containers.
- Amazon Elastic Container Service (Amazon ECS) orchestrates the running of Docker containers.
- Kubernetes is open-source software for container orchestration.
- Amazon Elastic Kubernetes Service (Amazon EKS) enables you to run Kubernetes on AWS
- Amazon Elastic Container Registry (Amazon ECR) enables you to store, manage, and deploy your Docker containers.

Building Careers
Through Education



Discussion

1. How does Azure handle Kubernetes?



Building Careers
Through Education

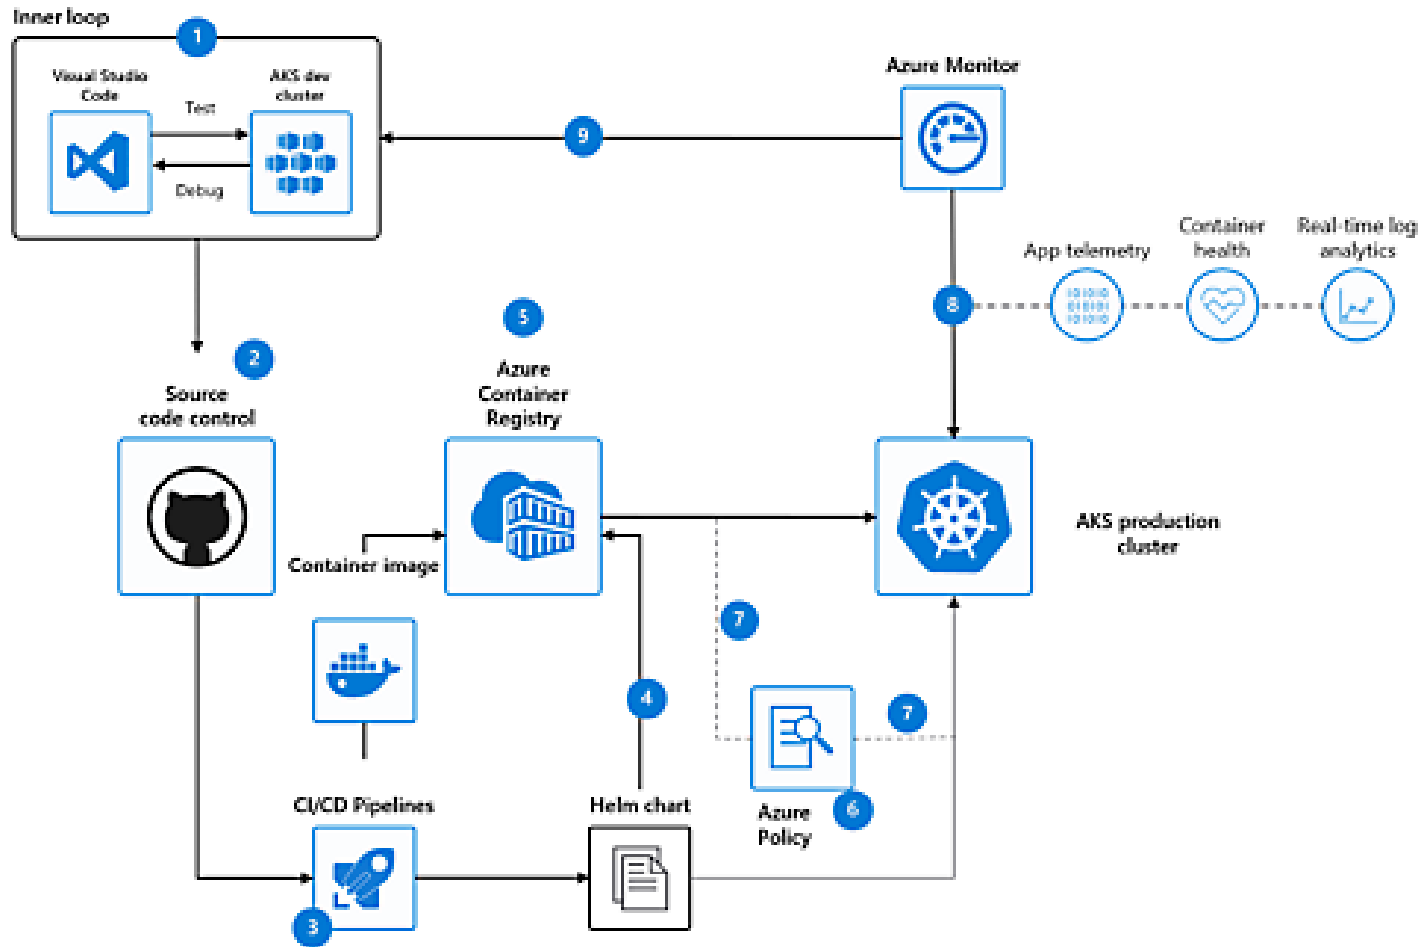


Submit your responses to
the chat!



Demo walkthrough

Kubernetes on Azure



Building Careers
Through Education







Demo walkthrough

Kubernetes on Azure

Home > Create a resource > Create Kubernetes cluster >
Choose a cluster preset configuration ...

Choose your scenario to view and apply the recommended configurations suited to your needs. The settings in the table below will be updated to the specified values based on your selection. All other cluster settings will remain unchanged. [Learn more](#)

	 Production Standard <input type="radio"/>	 Dev/Test <input checked="" type="radio"/>	 Production Economy <input type="radio"/>	 Production Enterprise <input type="radio"/>
	Best for most applications serving production traffic with AKS recommended best practices.	Best for developing new workloads or testing existing workloads.	Best for serving production traffic in a cost conscious way if your workloads can tolerate interruptions.	Best for serving production traffic with rigorous permissions and hardened security.
System node pool node size ⓘ	Standard_D8ds_v5 ⓘ	Standard_DS2_v2 ⓘ	Standard_D8ds_v5 ⓘ	Standard_D16ds_v5 ⓘ
System node pool autoscaling range ⓘ	2-5 nodes	2-100 nodes	2-5 nodes	2-5 nodes
User node pool node size ⓘ	Standard_D8ds_v5 ⓘ	-	Standard_D8as_v4 (SPOT instance) ⓘ	Standard_D8ds_v5 ⓘ
User node pool autoscaling range ⓘ	2-100 nodes	-	0-25 nodes	2-100 nodes
Private cluster ⓘ	-	-	-	✓
Availability zones ⓘ	✓	-	-	✓
Azure Policy ⓘ	✓	-	-	✓
Azure Monitor ⓘ	✓	-	-	✓
Secret store CSI driver ⓘ	-	-	-	✓
Network configuration ⓘ	Azure CNI Overlay ⓘ	Azure CNI Overlay ⓘ	Azure CNI Overlay ⓘ	Azure CNI Overlay ⓘ
Network policy ⓘ	None	None	None	None
Authentication and Authorization ⓘ	Local accounts with Kubernetes RBAC ⓘ	Local accounts with Kubernetes RBAC ⓘ	Microsoft Entra ID authentication with Azure RBAC ⓘ	Microsoft Entra ID authentication with Azure RBAC ⓘ

[Save](#) [Cancel](#)

Building Careers
Through Education




Demo walkthrough

Kubernetes on Azure

[Home](#) > [Create a resource](#) >



Create Kubernetes cluster ...

[Basics](#) [Node pools](#) [Networking](#) [Integrations](#) [Monitoring](#) [Advanced](#) [Tags](#) [Review + create](#)


Azure Kubernetes Service (AKS) manages your hosted Kubernetes environment, making it quick and easy to deploy and manage containerized applications without container orchestration expertise. It also eliminates the burden of ongoing operations and maintenance by provisioning, upgrading, and scaling resources on demand, without taking your applications offline. [Learn more](#) 


Project details


Select a subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription * 	<div>mySubscription</div>
Resource group * 	<div>(New) myResourceGroup</div> <div>Create new</div>

Cluster details


Cluster preset configuration *	<div> Dev/Test</div> <div>To quickly customize your Kubernetes cluster, choose one of the preset configurations above. You can modify these configurations at any time. Compare presets</div>
--------------------------------	--


Kubernetes cluster name * 	<div>myAKSCluster</div>
---	-------------------------

Region * 	<div>(US) East US 2</div>
--	---------------------------


Availability zones 	<div>None</div>
--	-----------------

AKS pricing tier 	<div>Free</div>
--	-----------------

Kubernetes version * 	<div>1.29.7 (default)</div>
--	-----------------------------


Automatic upgrade 	<div>Enabled with patch (recommended)</div>
---	---



Automatic upgrade scheduler	<div>Every week on Sunday (recommended)</div> <div>Start on: Sat Aug 10 2024 00:00 +00:00 (Coordinated Universal Time) Edit schedule</div>
-----------------------------	---

Node security channel type 	<div>Node Image</div>
--	-----------------------

Security channel scheduler	<div>Every week on Sunday (recommended)</div> <div>Start on: Sat Aug 10 2024 00:00 +00:00 (Coordinated Universal Time) Edit schedule</div>
----------------------------	---

Choose between local accounts or Microsoft Entra ID for authentication and Azure RBAC or Kubernetes RBAC for your authorization needs.

Authentication and Authorization 	<div>Local accounts with Kubernetes RBAC</div>
--	--


 Once the cluster is deployed, use the Kubernetes CLI to manage RBAC configurations. [Learn more](#) 

[Previous](#)

[Next](#)

[Review + create](#)

Building Careers Through Education



Demo walkthrough

Kubernetes on Azure

Add a node pool ...

myAKSCluster

Node pool name * ⓘ

Mode * ⓘ ☒ User ☐ System

OS SKU * ⓘ ☐ Azure Linux ☒ Ubuntu Linux ☐ Windows 2022 ☐ Windows 2019

Availability zones ⓘ

Enable Azure Spot instances ⓘ ☐

Node size * ⓘ **Standard D2s v3**
2 vcpus, 8 GiB memory
[Choose a size](#)

Scale method ⓘ ☐ Manual ☒ Autoscale - Recommended
ⓘ This option is recommended so that the cluster is automatically sized correctly for the current running workloads.

Minimum node count * ⓘ

Maximum node count * ⓘ
The maximum node count allowed for an AKS cluster is 1000 per node pool and 5000 nodes across all node pools in this cluster.

Optional settings

Max pods per node * ⓘ 10 - 250

Enable public IP per node ⓘ ☐


Labels
Labels are key/value pairs that can be used to categorize or add identifying information to Kubernetes resources such as nodes. Labels for the node pool will be applied to each node in the node pool. [Learn more ⓘ](#)

Key	Value
<input type="text"/>	<input type="text"/>

Taints
Taints are tuples that are used in conjunction with tolerations to determine which pods can be scheduled on which nodes. In order for a pod to be scheduled to a node, it must tolerate all of the taints applied to that node. Taints for the node pool will be applied to each node in the node pool. [Learn more ⓘ](#)

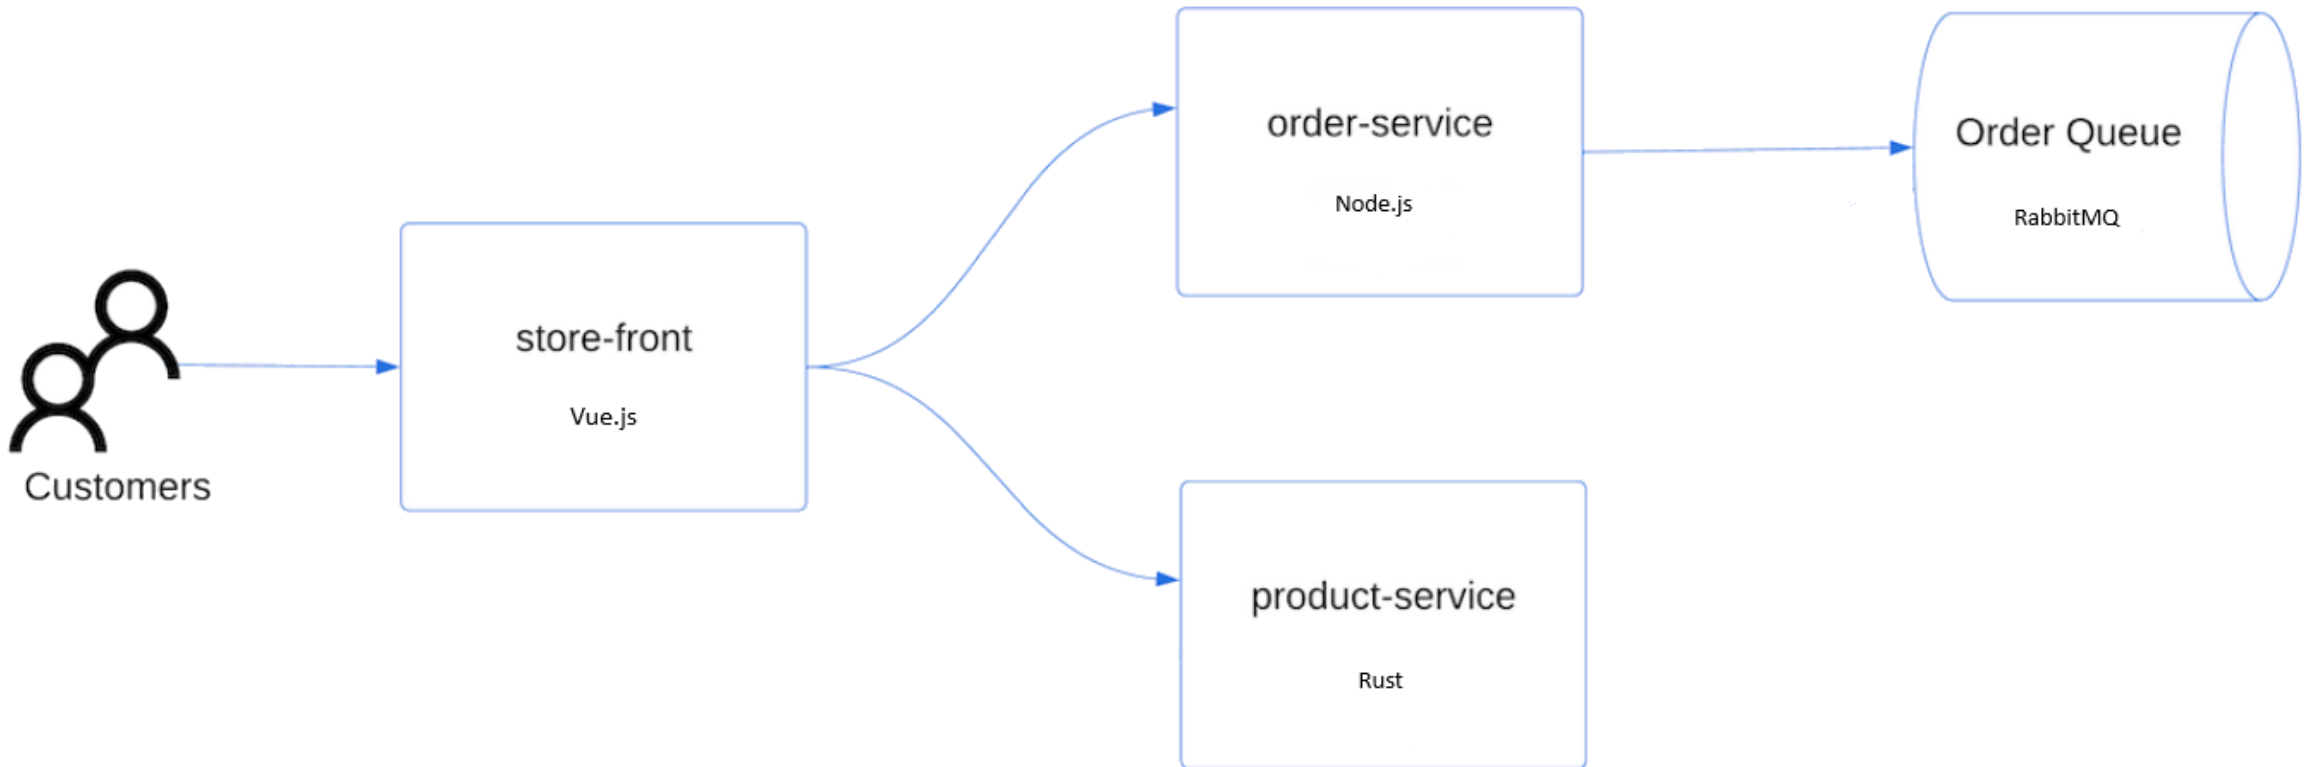
Key	Value	Effect
<input type="text"/>	<input type="text"/>	<input type="text"/>

Building Careers
Through Education

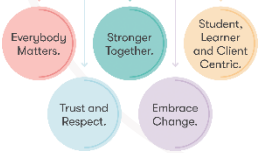


Demo walkthrough

Kubernetes on Azure




Building Careers
Through Education

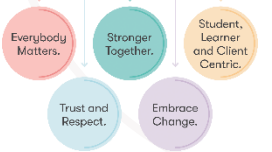


Demo walkthrough

The debrief



Building Careers
Through Education



Best practices for container orchestration

Use Health Checks:

- Define health checks to ensure containers are running as expected
- Automatically restart unhealthy containers

Implement Monitoring and Logging:

- Use tools like Prometheus, Grafana, and ELK stack for monitoring and logging
- Monitor resource usage and application performance


Optimise Resource Allocation:

- Define resource limits and requests for containers
- Use node affinity and anti-affinity rules to optimise placement

Ensure Security:

- Implement network policies and secrets management
- Regularly update and patch container images

Building Careers
Through Education




Demo walkthrough

Deploying a Python microservices application using Docker

Deploying a simple Python microservices application using Docker, Kubernetes and serverless deployments

- <https://www.eksworkshop.com/>
- <https://webapp.serverlessworkshops.io/>

Building Careers
Through Education



Post-webinar consolidation

Apply...

- **Task 1:** Deploy a sample microservices application using Docker and Kubernetes
- **Task 2:** Write a reflection on resource optimisation strategies for sustainable growth
- **Task 3:** Reflect on the benefits of containerisation in application deployment in your organisation


Building Careers
Through Education



Session objective review

- Discuss how cloud automation and orchestration can help streamline the management and maintenance of cloud-based systems and applications;
- Analyse how the use of tools and technologies such as Ansible, Puppet, and Chef may be used to automate routine tasks in order to reduce the potential for human error in cloud-based systems and applications;
- Use tools and technologies such as Kubernetes, Docker Swarm, and Mesos to manage containers and balance workloads to ensure they are working together effectively and high availability exists.

Building Careers
Through Education





Thank you

**Do you have any questions,
comments, or feedback?**

