

M7T4 - Monitoring an ingestion service and anomaly detection techniques

Lab Activity 1: Setting Up Prometheus and Grafana for Kafka Monitoring (1 hour)

Objective: Automate monitoring processes for data ingestion services using industry-standard tools.

Problem: Imagine your company, a data analytics firm, needs to monitor Kafka clusters to ensure efficient data ingestion and processing. Your task is to set up Prometheus and Grafana to monitor Kafka metrics.

Step-by-step instructions

Setup:

Install Prometheus:

- **Step 1:** Download the latest Prometheus release from the official website.
- **Step 2:** Extract the downloaded file and navigate to the Prometheus directory.
- **Step 3:** Configure the prometheus.yml file to scrape metrics from

Kafka:global:

scrape_interval: 15s

scrape_configs:

- job_name: 'kafka_brokers'

static_configs:

- targets: ['localhost:9090']

- job_name: 'kafka_exporter'

static_configs:

- targets: ['localhost:9308']

- **Step 4:** Start Prometheus by running the following command: `./prometheus --config.file=prometheus.yml`

Install Grafana:

- **Step 1:** Download Grafana from the official website.
- **Step 2:** Install Grafana by following the installation instructions for your operating system.

- **Step 3:** Start Grafana by running the following command: `sudo service grafana-server start`

Configure Grafana:

- **Step 1:** Open your web browser and navigate to <http://localhost:3000>.
- **Step 2:** Log in with the default credentials (username: admin, password: admin).
- **Step 3:** Add Prometheus as a data source:
 - Go to **Configuration > Data Sources > Add data source**.
 - Select **Prometheus** and enter the URL <http://localhost:9090>.
 - Click **Save & Test**.

Create Dashboards:

- **Step 1:** Go to **Create > Dashboard**.
- **Step 2:** Add a new panel and select the metrics you want to visualise (e.g., kafka_brokers, kafka_exporter).
- **Step 3:** Customise the dashboard layout and save it.

Set Up Kafka Exporter:

- **Step 1:** Download the Kafka Exporter binary from the official repository.
- **Step 2:** Run the Kafka Exporter with the following command: `./kafka_exporter -kafka.server=localhost:9092`

Monitor Key Metrics:

- **Step 1:** In Grafana, create panels to monitor key Kafka metrics such as under-replicated partitions, offline partitions, active controller count, request latency, error rates, consumer lag, and consumption rate.
- **Step 2:** Set up alerts for critical metrics using Prometheus Alertmanager.

Expected Outcomes:

- **Prometheus Setup:** Prometheus is installed and configured to scrape Kafka metrics.
- **Grafana Setup:** Grafana is installed and configured to visualise Kafka metrics.
- **Kafka Exporter Setup:** Kafka Exporter is running and providing metrics to Prometheus.
- **Dashboard Creation:** Grafana dashboards are created to monitor key Kafka metrics.
- **Alert Configuration:** Alerts are set up for critical Kafka metrics.

Sample Data:

- **Before Monitoring:** global:
scrape_interval: 15s

scrape_configs:
- job_name: 'kafka_brokers'

```
static_configs:
  - targets: ['localhost:9090']
```

```
- job_name: 'kafka_exporter'
  static_configs:
    - targets: ['localhost:9308']
```

- **After Monitoring:** !Grafana Dashboard

Lab Activity 2: Implementing Anomaly Detection with Isolation Forest (1 hour)

Objective: Implement forecasting and anomaly detection techniques, including ARIMA, SARIMAX, and other methods.

Problem: Imagine your company, a cybersecurity firm, needs to detect anomalies in network traffic data to identify potential security threats. Your task is to implement anomaly detection using the Isolation Forest algorithm.

Step-by-step instructions

Setup:

Set Up the Environment:

- **Step 1:** Create a new Python virtual environment:

```
python3 -m venv anomaly_detection_env
source anomaly_detection_env/bin/activate
```
- **Step 2:** Install the required libraries:

```
pip install pandas scikit-learn matplotlib
```

Load and Prepare Data:

- **Step 1:** Create a sample dataset with time series data:

```
import pandas as pd
import numpy as np
```

```
# Generate sample data
np.random.seed(42)
dates = pd.date_range('20230101', periods=100)
data = pd.DataFrame(np.random.randn(100, 1), index=dates,
columns=['value'])
data['value'] += np.linspace(0, 10, 100)
data['anomaly'] = 0
data.loc[50, 'value'] = 20 # Inject an anomaly
data.loc[50, 'anomaly'] = 1
```

Train Isolation Forest Model:

- **Step 1:** Import the Isolation Forest model and train it on the dataset:from
sklearn.ensemble import IsolationForest

```
# Train Isolation Forest model
model = IsolationForest(contamination=0.01, random_state=42)
model.fit(data[['value']])
data['anomaly_score'] = model.decision_function(data[['value']])
data['anomaly'] = model.predict(data[['value']])
data['anomaly'] = data['anomaly'].apply(lambda x: 1 if x == -1 else 0)
```

Visualise Anomalies:

- **Step 1:** Plot the data and highlight the detected anomalies:import
matplotlib.pyplot as plt

```
plt.figure(figsize=(10, 6))
plt.plot(data.index, data['value'], label='Value')
plt.scatter(data.index, data['value'], c=data['anomaly'], cmap='coolwarm',
            label='Anomaly')
plt.title('Anomaly Detection with Isolation Forest')
plt.xlabel('Date')
plt.ylabel('Value')
plt.legend()
plt.show()
```

Interpret Results:

- **Step 1:** Discuss the results and the effectiveness of the Isolation Forest model in detecting anomalies.
- **Step 2:** Explore how different contamination levels and parameters affect the model's performance.

Expected Outcomes:

- **Environment Setup:** Python virtual environment is created and required libraries are installed.
- **Data Preparation:** Sample dataset with time series data is created and anomalies are injected.
- **Model Training:** Isolation Forest model is trained on the dataset.
- **Anomaly Detection:** Anomalies are detected and highlighted in the dataset.
- **Result Interpretation:** Results are discussed, and the model's performance is evaluated.

Sample Data:

- **Before Anomaly Detection:**dates = pd.date_range('20230101', periods=100)
data = pd.DataFrame(np.random.randn(100, 1), index=dates,
columns=['value'])
data['value'] += np.linspace(0, 10, 100)

```
data['anomaly'] = 0
data.loc[50, 'value'] = 20 # Inject an anomaly
data.loc[50, 'anomaly'] = 1
```

- **After Anomaly Detection:**

```
plt.figure(figsize=(10, 6))
plt.plot(data.index, data['value'], label='Value')
plt.scatter(data.index, data['value'], c=data['anomaly'], cmap='coolwarm',
            label='Anomaly')
plt.title('Anomaly Detection with Isolation Forest')
plt.xlabel('Date')
plt.ylabel('Value')
plt.legend()
plt.show()
```