




Data engineering

Module: Data pipelines

Topic: From data pipelines to data architectures

**Welcome to today's
webinar.**



Ice breaker

Scale it or fail it!

“Imagine you’re hosting a dinner party for 4 people. Suddenly, 40 guests show up. What’s your first move?”

Your options:

- A. Order 10 pizzas and hope for the best
- B. Start cooking in batches and delegate tasks
- C. Set up a buffet line and let people serve themselves
- D. Panic and pretend it was a prank

Building Careers
Through Education



**Submit your responses to the
chat or turn on your
microphone**



Real-world case study

Zalando's scalable data platform on AWS

Challenge

- Billions of daily events from web, mobile, and logistics
- Real-time personalization and inventory updates
- Elastic scaling for peak events like Black Friday

Solution

- Kafka + Kinesis for real-time streaming
- Serverless compute with AWS Lambda
- Modular microservices and Airflow orchestration
- Scalable storage with S3 and Redshift

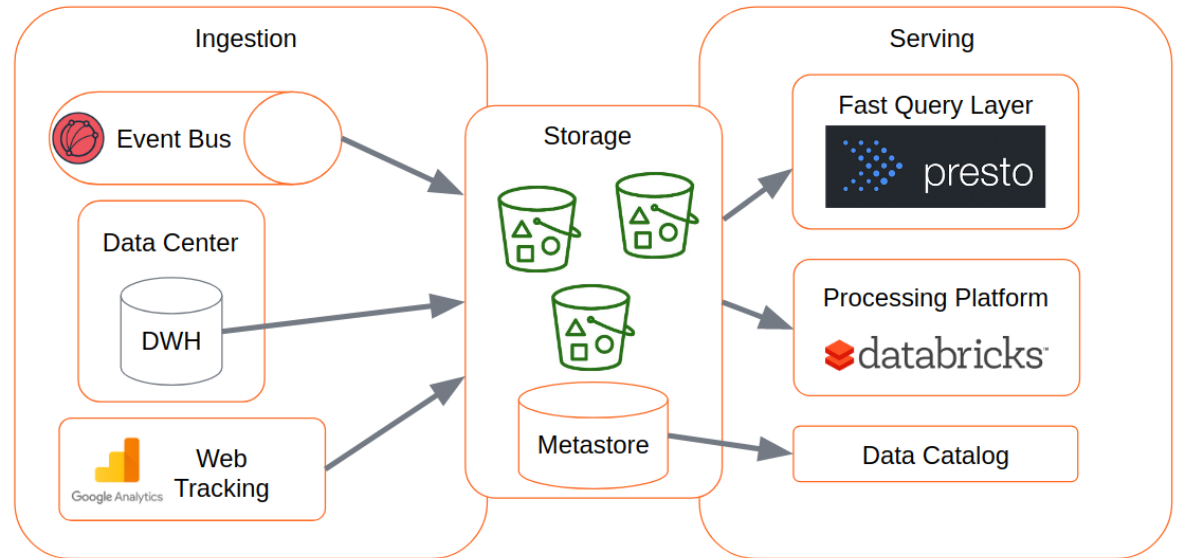


Zalando delivers real-time recommendations, accurate inventory, rapid feature rollouts, and seamless performance during traffic spikes.

Building Careers
Through Education



Zalando's Data Lake



Zalando's Real-Time Retail Engine: Scaling with AWS,
Image source: [AWS](#)



e-learning recap

Reflecting on your learning...

The e-learning for this topic, covered the following areas:

1. Scalability means designing systems that grow with data and demand
2. Pipelines scale through horizontal scaling, stream processing, and decoupling
3. Performance and reliability require monitoring, fault tolerance, and testing
4. Scalable architectures are modular, loosely coupled, and event-driven
5. Tools like Spark, Kafka, and Airflow enable distributed, resilient workflows
6. Smart trade-offs balance speed, cost, and complexity



- Do you have any questions about any of these areas?
- Did everything in the e-learning make sense?

Building Careers
Through Education



Q&A discussion




Webinar Agenda

Today, we will cover the following:

1. Core concept recap
2. Practical lab
3. Summary
4. Q&A

Building Careers
Through Education




Session aim and objectives

By the end of this session, you should be able to:


- Understand what scalability means in data pipelines and architectures
- Explore strategies and tools for building scalable systems
- Learn from real-world examples of scalable data design

Building Careers
Through Education






Recap of core concepts



What is scalability?

From prototype to production

- Scalability is the ability of a system to grow with demand
- It applies to both data pipelines and overall architecture
- Critical for performance, reliability, and cost-efficiency



Scalability is the ability of your data solution to grow as demand increases.




Figure: An abstraction of a data pipeline

Building Careers
Through Education



Scaling data pipelines

Strategies for growth





Figure: Vertical vs. Horizontal Scaling



Batch pipeline

Like a postal truck doing pickups and deliveries once a day.

Vs



Streaming pipeline

Like a high-speed conveyor belt: always on, always moving.

Figure: Difference between Batch and Streaming via an analogy.




Figure: Diagram of how MapReduce works

Building Careers
Through Education



Tools for scalable pipelines

The Engineer's Toolkit



Apache Spark: distributed processing for large-scale data



Apache Kafka: real-time messaging and decoupling



Airflow & Prefect: orchestrate complex workflows


Building Careers
Through Education



Ensuring performance

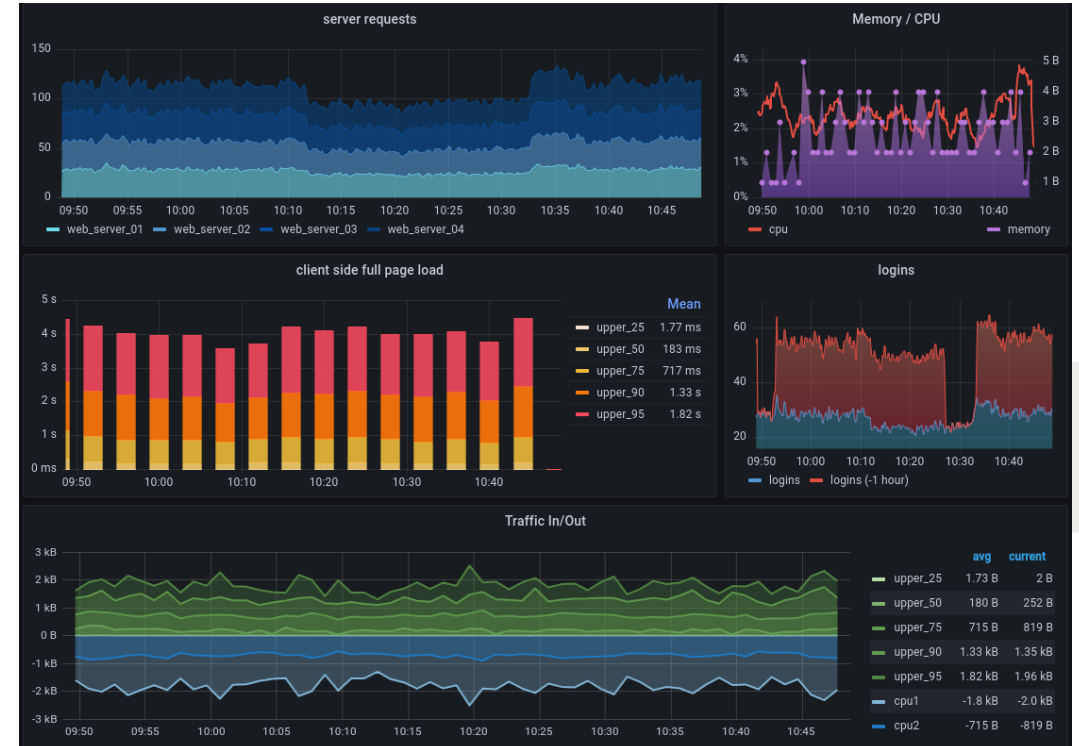
Monitoring and optimisation

- Track throughput, latency, and resource usage
- Optimise transformations to reduce compute load
- Choose storage formats that match your workload



Scalable systems depend on visibility - monitoring and optimising performance is essential to ensure your pipelines stay fast, efficient, and reliable as they grow.

Building Careers
Through Education




Ensuring Performance: Monitoring and optimisation strategies for scalable data pipelines, Image source: [Grafana.com](https://grafana.com)



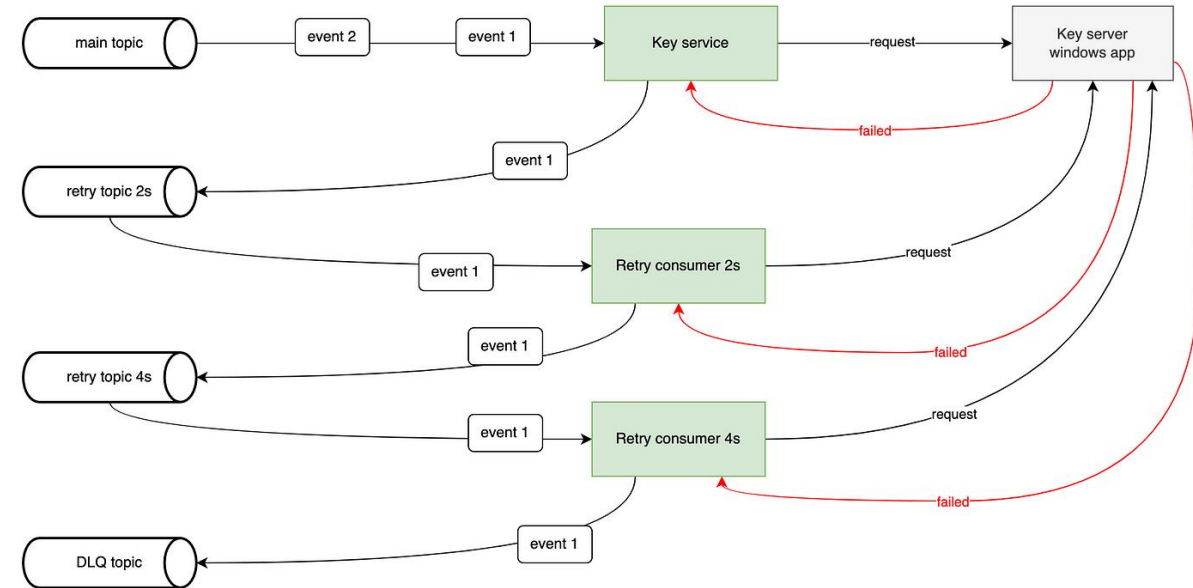
Building for reliability

Fault tolerance and resilience

- Use retries, checkpoints, and idempotent design
- Decouple components with queues like Kafka
- Validate data early and test beyond unit tests



Reliability isn't about avoiding failure - it's about designing systems that recover gracefully, validate early, and keep data trustworthy at scale.




Resilient Event Processing: How Retry Logic and Dead Letter Queues Keep Pipelines Reliable, image source: [Medium.com](https://medium.com)




Designing scalable architectures

Thinking beyond pipelines

- Modular, loosely coupled components scale better
- Event-driven design enables real-time responsiveness
- Use the right storage for the right data type



Scalable architectures are built, not bolted on - modularity, event-driven design, and smart storage choices are the foundation for systems that grow and evolve.



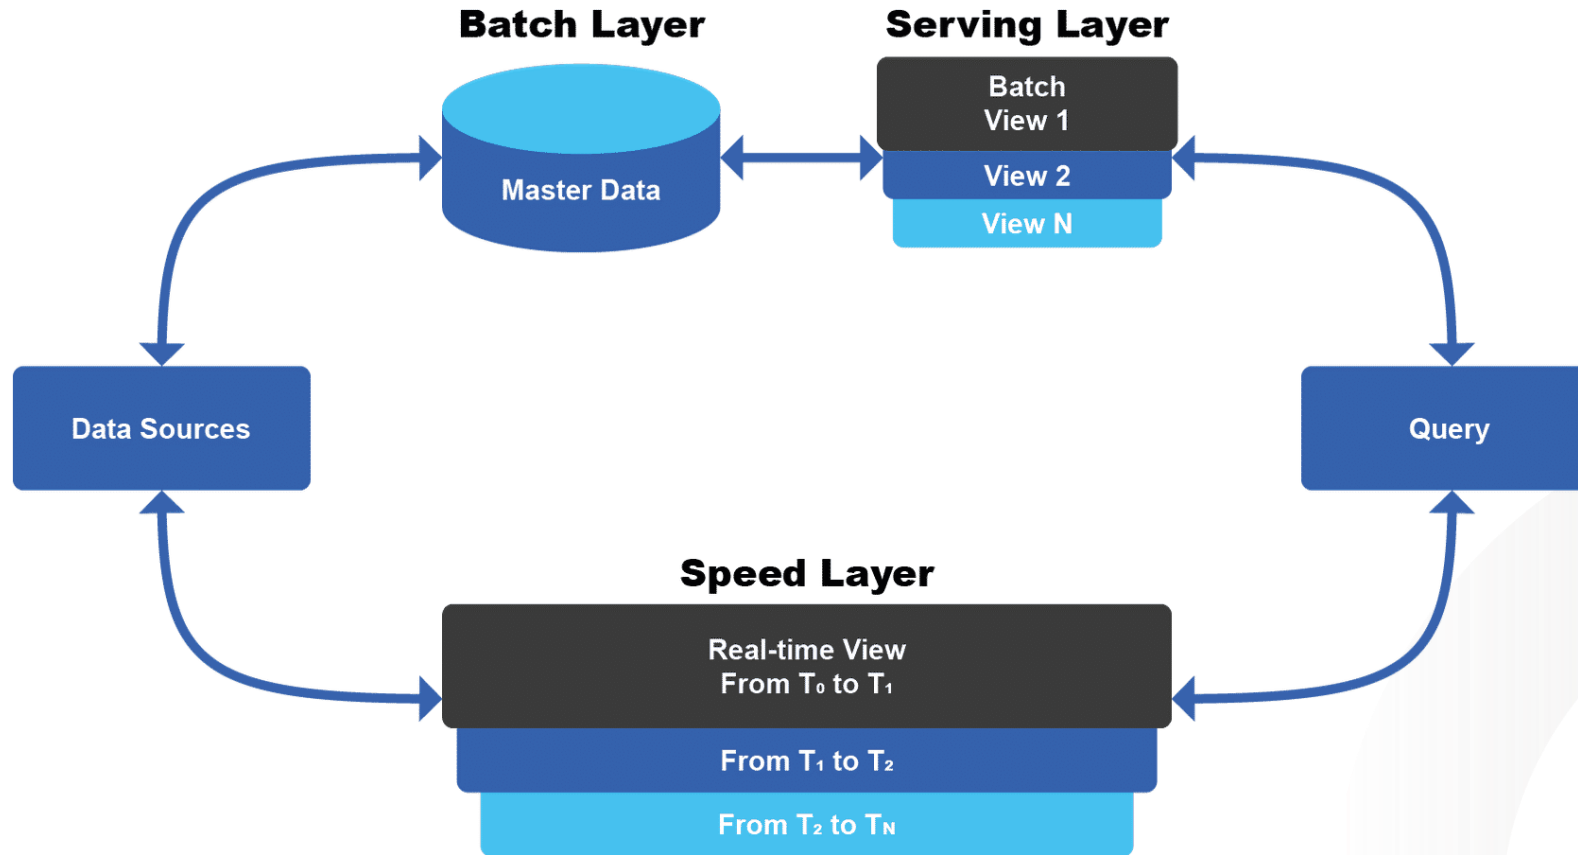
Scalable by Design: Modular, Architectures That Grow with Your Data,
image source: [freepic.com](https://www.freepic.com)

Building Careers
Through Education



Architecture patterns

Lambda



Lambda Architecture combining batching and streaming

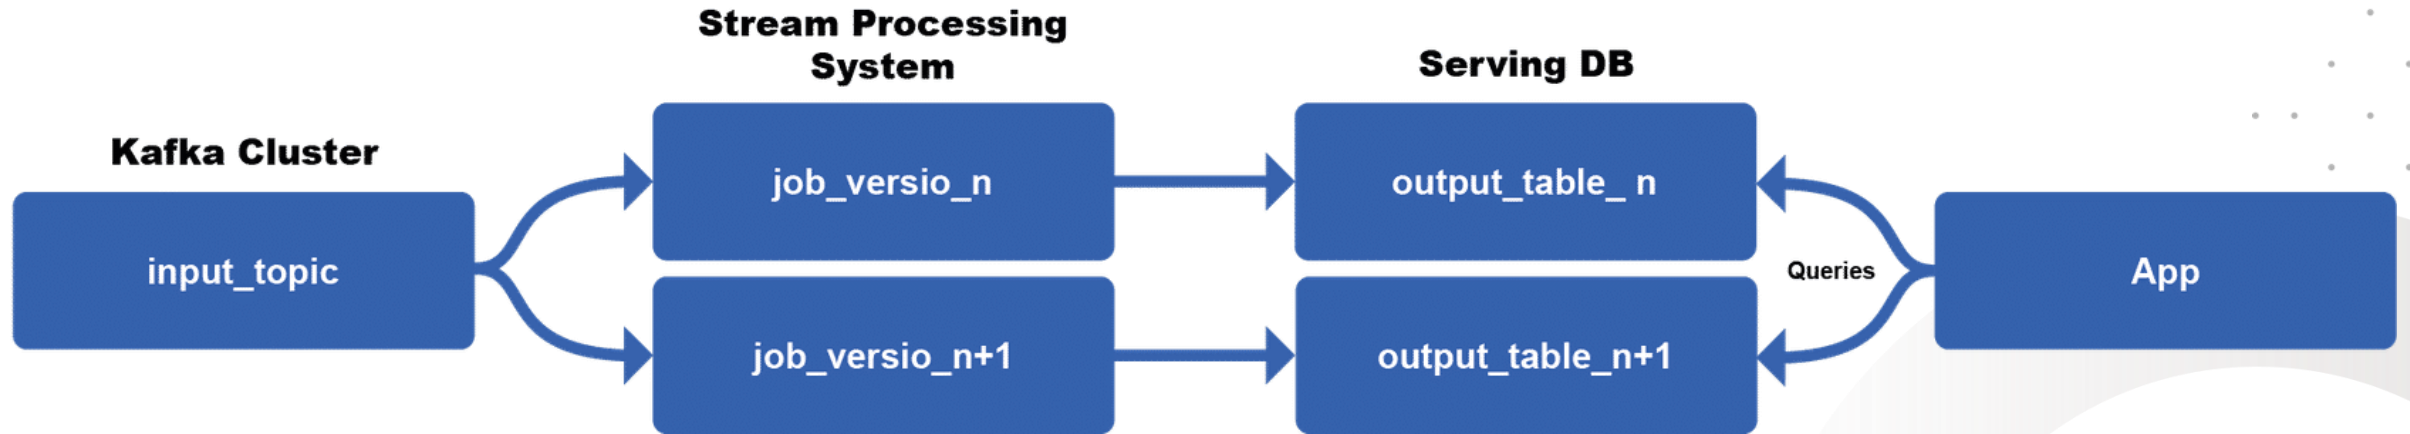
Building Careers
Through Education



Architecture patterns

Kappa

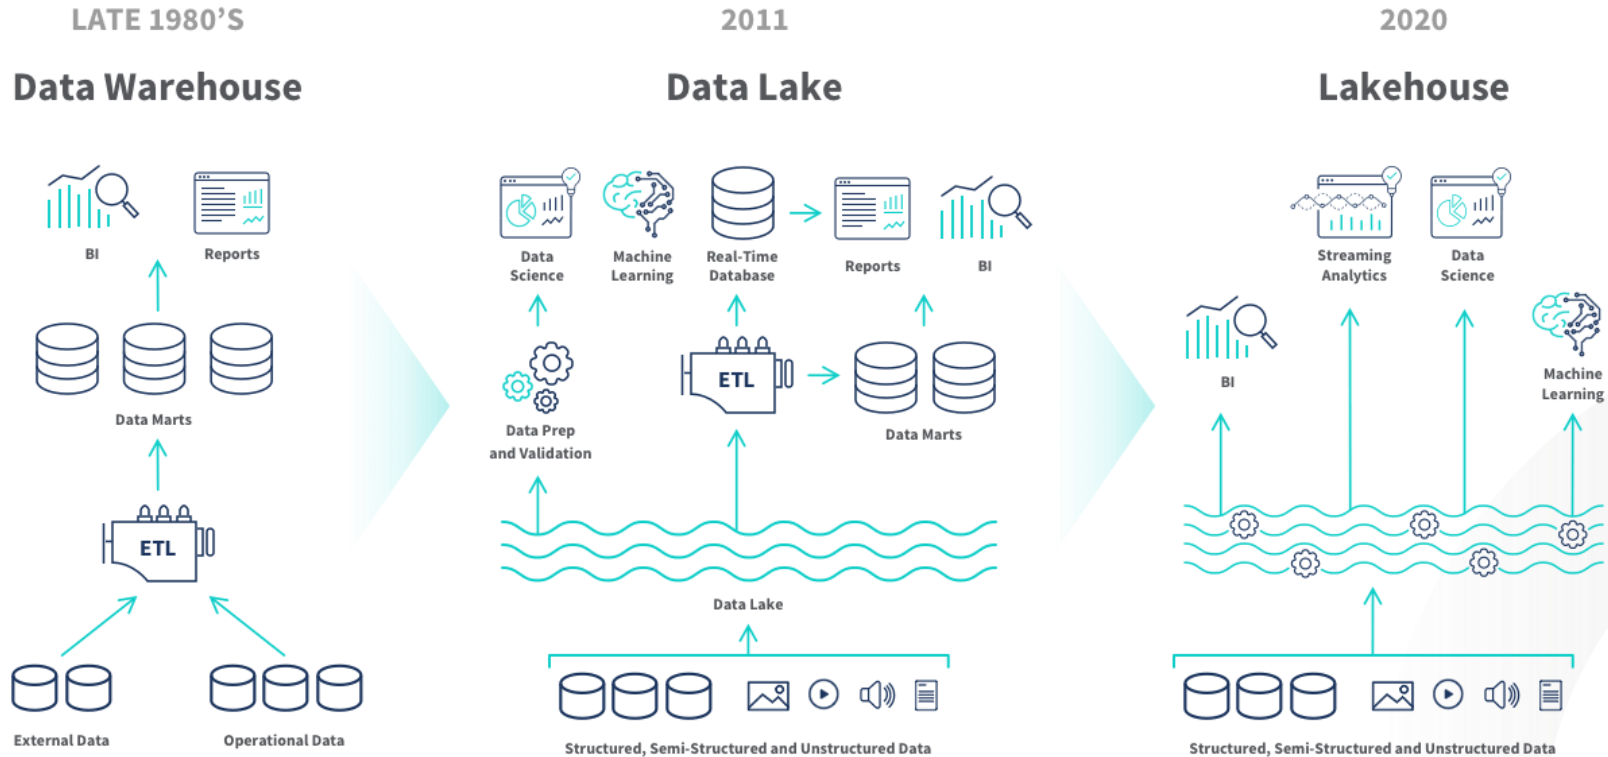
Building Careers
Through Education



Kappa Architecture for streaming based process

Architecture patterns

Lakehouse




Kappa Architecture, which is a purely streaming based process, image source: [HAZELCAST](#)



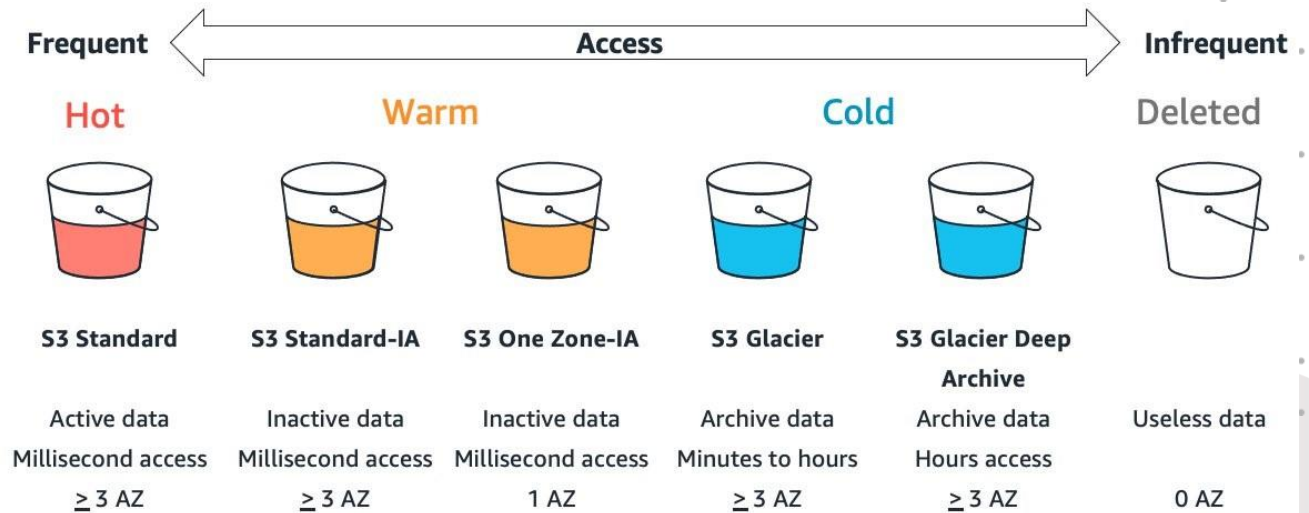
Balancing cost and performance

Smart scaling

- Store cold data cheaply; query hot data efficiently
- Use batch jobs to reduce always-on compute costs
- Auto-scale with guardrails to avoid surprise bills



Scalability isn't just about speed - it's about making smart, cost-effective choices that align with your workload and business goals.



Smart Scaling in the Cloud: Matching Storage and Compute to Your Workload Needs

Real – World case studies

Lessons from the field...

Building Careers
Through Education




Ride-sharing: Kafka + Flink for real-time GPS data

Image source: ottocar.co.uk



Retail: Lakehouse model for unified analytics

Image source: Webandcrafts.com




Media: Auto-scaling clusters for burst traffic

Image source: BBC.co.uk



Practical lab



Exercise part 4


Next phase of design based on Scottish data

- **Update the Database Schema:** Update your database schema to accommodate data from the Scottish version of the software application.
- **Re-implement the Data Pipeline:** Re-implement the pipeline to clean, validate, and load the four data sets into your revised database schema.
- **Utilise Additional Test Data:** Use additional test data provided by the Scottish application team, including data for ten Scottish users and the login timestamps.
- **Stakeholder Support:** Your instructor will act as the stakeholder SME and answer any questions about the new data.
- **Documentation Requirements:** Create full documentation for both the schema and the pipeline for implementation by other teams.

Files Provided in the Hub:

- The 10 sample records are in the "SC User Data.csv" file
- The login audit file is "SC-User-LoginTS.csv."

Building Careers
Through Education



Practical challenge




Key Learning Summary

- **Scalability Fundamentals** - Systems must handle growing data volume, complexity, and user demand without performance loss—across both pipelines and architectures.
- **Scaling Data Pipelines** - Use horizontal scaling, stream processing, partitioning, and decoupling (e.g., Kafka) to improve throughput and resilience.
- **Key Tools for Scale** - Leverage Apache Spark for distributed processing, Kafka for real-time streaming, and Airflow/Prefect for orchestration.
- **Performance & Reliability** - Monitor metrics like latency and resource usage, optimise transformations, and build in retries, validation, and fault tolerance.
- **Architectural Design Principles** - Design modular, loosely coupled systems using event-driven flows and polyglot persistence for flexibility and maintainability.
- **Architecture Patterns** - Choose from Lambda (batch + stream), Kappa (stream-only), or Lakehouse (unified analytics) based on your data needs.
- **Cost-Performance Trade-offs** - Store cold data cheaply, use batch jobs to reduce compute costs, and auto-scale with guardrails to avoid surprise bills.



Any questions or feedback?



Building Careers
Through Education





Thank you

